

- Define and architect AI agents for Supply Chain use cases, including frameworks for agent orchestration, multi-agent coordination, RAG, deployment, monitoring, and life cycle management.
 - Be hands on in quick proof of concepts development to demonstrate technical feasibility and implement enterprise grade Agentic Supply Chain solutions
 - Partner with Enterprise IT engineering, product, and research teams while evaluating LLMs, agentic frameworks, and NVIDIA's own NeMo technologies
 - Ensure integration with enterprise IT and Operations data sources and Industry's best Agentic platforms with strong content security focus.
 - Drive architectural decisions across deployment models (on-prem, cloud, hybrid, containerized) to deliver scalable, reliable, and efficient solutions.
 - Lead design reviews, develop technical documentation, and mentor engineers, data scientists and solution architects in principles of architecture and code development.
 - Champion observability, monitoring, versioning, and telemetry to ensure trustworthy and auditable AI agents.
 - Influence Supply Chain Operations adoption of the platform by partnering with stakeholders across IT, supply chain and serve as a reference adopter providing feedback to strengthen NVIDIA's ecosystem.
-
- Define the AI Vision for Commerce: Translate creative and transactional intent into agentic workflows that span discovery, personalization, and conversion.
 - Lead Roadmap Strategy: Collaborate with product and engineering leaders to prioritize AI capabilities, including LLM-powered assistants, retrieval-augmented generation (RAG), and autonomous agents.
 - Prototype and Evangelize: Build proof-of-concepts that demonstrate the power of agentic systems in commerce. Evangelize internally and externally through demos, talks, and documentation.
 - Guide Technical Implementation: Partner with engineering teams to integrate AI services into Adobe Commerce, ensuring scalability, performance, and ethical alignment.
 - Shape Data Strategy: Influence how semantic data is captured, embedded, and used to train and fine-tune models for commerce-specific use cases.
 - Monitor and Optimize: Track model performance post-deployment, recommend improvements, and ensure continuous learning and adaptation.
 - Be the Voice of AI: Represent Adobe Commerce's AI strategy in cross-functional forums, customer engagements, and industry conversations.

What You Bring:

- 10+ years of experience in software architecture, AI/ML systems, or product innovation.
- Deep expertise in LLMs, RAG, agentic systems, and AI orchestration frameworks.
- Hands-on experience deploying AI/ML solutions in production environments.

- Strong coding skills in JavaScript, React, GraphQL, and at least one backend language (Java, Kotlin, Python).
 - Familiarity with microservices, cloud-native architectures, and DevOps practices.
 - Proven ability to influence senior stakeholders and drive cross-functional alignment.
 - Passion for customer-centric design and ethical AI development.
 - Undergraduate degree required; graduate degree in a quantitative field pre
-
- Serve as an expert in implementing and deploying AI applications based on large language models (LLMs), internal and external agentic frameworks, and custom models.
 - Work with hardware architects to identify how to best design, customize, and deploy AI-based solutions to their specific problem domains.
 - Collaborate with infrastructure engineers to improve existing automated workflows by incorporating LLMs and establishing best practices for future solutions.
 - Develop and optimize retrieval and generation algorithms for enterprise data (text, code, and images) to build advanced AI applications.
 - Interact with internal research groups on how to solve complex chip design problems in new ways by leveraging machine learning (ML) and deep learning (DL).
 - Research emerging AI technologies and engineering best practices to continuously evolve our development ecosystem and maintain a competitive edge.

What We Need To See

- MSc or PhD in Data Science, Computer Science/Engineering, Electrical Engineering, or equivalent experience.
 - 5+ years of industry or research experience.
 - Deep practical knowledge of LLMs, DL/ML, and Agent development.
 - Well versed in agentic literature and eager to continue learning.
 - Strong background in implementing AI solutions to solve real-world engineering problems.
 - Experience with training/fine-tuning custom models, building multi-agent systems, retrieval augmented generation (RAG) pipelines, and vector databases.
 - Strong analytical, communication, and interpersonal skills.
-
- Define and maintain the overarching GenAI, LLMs, RAG pipelines, and autonomous agent systems
 - Design and implement multiple agent orchestration workflows and agentic frameworks.
 - Evaluate and select AI/Agentic tools, frameworks, and platforms (ex: LangChain, Semantic Kernel, Vertex AI, Azure OpenAI, AWS Bedrock, LangGraph, CrewAI)
 - Align architectural decisions with business performance metrics, latency, security, cost, and explainability requirements.
 - Assess emerging AI technologies and trends, recommending their adoption where appropriate

- Design scalable architectures that support composability, modularity, observability, and scalability of AI solutions, ensuring alignment with business strategy, data strategy, and technology roadmap
- Participate as a key stakeholder in the development of our AI Roadmap

Solution Design:

- Collaborate with cross-functional teams, including data architects, data scientists, enterprise architecture, security, AI product managers, and business stakeholders
- Guide proof-of-concept development and prototype evaluations to validate architecture decisions
- Provide architecture review, feedback, and mentorship to AI engineering teams
- Conduct structured build-vs-buy evaluations and contribute to platform roadmap decisions

Governance:

- Embed responsible AI principles, privacy, and governance frameworks into system design
- Ensure AI architecture complies with enterprise standards, security policies, and regulatory frameworks
- Partner with AI governance, data governance, security, and compliance teams to implement transparency and auditability

Qualifications

Technical Experience:

- BA/BS plus at least 10 years of relevant architecture experience or demonstrated equivalency of experience and/or education
- At least 4 years of experience in ML or AI systems design and architecture
- Proven expertise in designing and deploying enterprise-grade genAI, RAG models, agentic AI, and ML pipelines.
- Working knowledge/experience with interoperability protocols such as Model Context Protocol (MCP) and Agent-to-Agent (A2A) for cross-platform agent communication.
- Experience designing/implementing solutions with agentic frameworks (e.g., LangChain, Azure AI Agent Services, n8n, Autogen, etc.).
- Working knowledge of AI Gateway implementation for enforcing guardrails, monitoring, and centralized model access control.
- At least 3 years of experience building AI solutions in AWS or equivalent
- Ability to build AI POCs and design for enterprise production

Leadership & Collaboration:

- Excellent collaboration and communication skills are key
- Must be comfortable engaging with key stakeholders and IT top-level leadership
- Proven ability to lead and influence cross-functional teams, including software engineer teams and software product managers
- Demonstrated ability to stay on top of AI trends and best practices

Additional Information

Preferred Qualifications

- Understanding of MLOps and DevOps practices including CI/CD for models, observability, and rollback GitHub, GitHub Actions

- CI/CD pipelines using YML, AWS CloudFormation, Terraform
 - Is intellectually curious in order to keep on top of new concepts in AI/ML, trends, and best practices
 - Containerization tools like Kubernetes and Docker
 - Programming languages/frameworks including C#, Python, JavaScript, JSON
 - Azure DevOps for recording and tracking of Epics, Features and User Stories
-
- Architect and deliver enterprise-scale AI/ML solutions including data pipelines, feature engineering, model development, deployment, monitoring, and lifecycle management.
 - Design scalable, secure architectures that integrate across cloud-native and hybrid environments (AWS, Azure, GCP).
 - Lead projects across applied AI domains such as natural language processing, computer vision, time series forecasting, recommender systems, and generative AI.
 - Enable GenAI and LLMOps adoption, including fine-tuning, RAG (retrieval-augmented generation), prompt engineering, and orchestration frameworks.
 - Champion Responsible AI practices (explainability, fairness, privacy, bias mitigation, and governance) to ensure ethical and compliant AI delivery.
 - Evaluate emerging technologies, frameworks, and tools, providing clear recommendations tailored to client use cases.
 - Lead and mentor cross-functional delivery teams of data engineers, ML engineers, and consultants, ensuring on-time, high-quality outcomes.
 - Support business development through proposals, client pitches, solution accelerators, and thought leadership.
 - Coach and mentor junior consultants, fostering a culture of continuous learning and innovation across the AI/ML practice.

What You'll Bring

- 6+ years of experience implementing ML/AI models in production (classical ML, deep learning, or generative AI).
- 3+ years of experience in professional consulting or IT services, with proven ability to lead client-facing engagements.
- Deep expertise in ML/DL (TensorFlow, PyTorch, Scikit-learn).
- Demonstrated proficiency in GenAI frameworks (Hugging Face, LangChain).
- Strong programming skills in Python, R, or Scala with hands-on experience building scalable ML pipelines.
- Proficiency in cloud AI/ML services (AWS SageMaker, Azure ML, GCP Vertex AI) and big data ecosystems (Spark, Databricks, Kafka, Snowflake).
- Experience with MLOps/LLMOps, CI/CD, observability, and containerization (Docker, Kubernetes).
- Strong understanding of AI solution architecture, balancing performance, scalability, cost, and governance.

- Ability to communicate complex AI concepts to both technical and non-technical stakeholders, influencing executive decision-making.
 - Experience managing delivery teams and shaping AI/ML roadmaps for enterprise clients.
 - Strong problem-solving, critical thinking, and business acumen, with an innovative mindset.
-
- Architect and deliver enterprise-scale AI/ML solutions including data pipelines, feature engineering, model development, deployment, monitoring, and lifecycle management.
 - Design scalable, secure architectures that integrate across cloud-native and hybrid environments (AWS, Azure, GCP).
 - Lead projects across applied AI domains such as natural language processing, computer vision, time series forecasting, recommender systems, and generative AI.
 - Enable GenAI and LLMOps adoption, including fine-tuning, RAG (retrieval-augmented generation), prompt engineering, and orchestration frameworks.
 - Champion Responsible AI practices (explainability, fairness, privacy, bias mitigation, and governance) to ensure ethical and compliant AI delivery.
 - Evaluate emerging technologies, frameworks, and tools, providing clear recommendations tailored to client use cases.
 - Lead and mentor cross-functional delivery teams of data engineers, ML engineers, and consultants, ensuring on-time, high-quality outcomes.
 - Support business development through proposals, client pitches, solution accelerators, and thought leadership.
 - Coach and mentor junior consultants, fostering a culture of continuous learning and innovation across the AI/ML practice.

What You'll Bring

- 6+ years of experience implementing ML/AI models in production (classical ML, deep learning, or generative AI).
- 3+ years of experience in professional consulting or IT services, with proven ability to lead client-facing engagements.
- Deep expertise in ML/DL (TensorFlow, PyTorch, Scikit-learn).
- Demonstrated proficiency in GenAI frameworks (Hugging Face, LangChain).
- Strong programming skills in Python, R, or Scala with hands-on experience building scalable ML pipelines.
- Proficiency in cloud AI/ML services (AWS SageMaker, Azure ML, GCP Vertex AI) and big data ecosystems (Spark, Databricks, Kafka, Snowflake).
- Experience with MLOps/LLMOps, CI/CD, observability, and containerization (Docker, Kubernetes).
- Strong understanding of AI solution architecture, balancing performance, scalability, cost, and governance.
- Ability to communicate complex AI concepts to both technical and non-technical stakeholders, influencing executive decision-making.
- Experience managing delivery teams and shaping AI/ML roadmaps for enterprise clients.

- Strong problem-solving, critical thinking, and business acumen, with an innovative mindset.

Primary Responsibilities

- Involved across many aspects of architecture, design, build, engineering and implementation of technical solutions and capabilities providing technical leadership
- Communicate overall architecture vision, technical strategies, and trade-offs across various levels in the organization to gain buy-in
- Provide recommendations and technical guidance to improve performance, reliability, and reusability within the constraints of budget, resource and business dependencies
- Engage across teams in a capacity that ranges from assisting on in-flight initiatives, up to technical briefings and demonstrations of new technologies across the organization
- Leads business solutions, capabilities, and is a solid influencer that engages with client executives and stakeholders
- Provides leadership to and is accountable for the performance and results through multiple layers of management and senior level professional staff

- 8+ years of software engineering experience with Java and/or Python
- 5+ years of knowledge in generative AI models, AgenticAI systems for autonomous decision-making, and traditional AI/ML techniques like supervised, unsupervised, and reinforcement learning
- 3+ years of design and implementation of advanced AI/ML strategies
- 3+ years of Cloud-Native AI deployment
- 2+ years of experience with frameworks such as Langgraph, Autogen, CrewAI, OpenAI APIs, and Scikit-learn
- 1+ years of experience developing applications powered by generative AI (e.g., conversational AI, content creation, pers

5+ years of experience in AI/ML architecture with 8+ years in enterprise AI solutions

Deep expertise in LLM architectures, prompt engineering, and agentic frameworks (LangGraph, LangMem)

Hands-on experience with Azure OpenAI GPT-4/5, embedding models, and Azure cloud services

Strong background in Python, distributed systems, and enterprise architecture

Experience with Claude Code for agentic coding and AI-powered development

Proven track record in financial services or regulatory compliance environments

Expert knowledge of RAG architectures, advanced RAG patterns, and vector database optimization

Experience with Small Language Models (SLM), Agent-to-Agent (A2A) communication, and Model Context Protocol (MCP)

Proven ability to architect and scale AI solutions for enterprise workloads (1M+ documents, sub-second response times)

Key Responsibilities

Design end-to-end AI solutions for private equity fund operations and financial automation

Architect scalable agentic AI frameworks using LangGraph, LangMem, and custom agent orchestration

Lead technical strategy for Azure OpenAI GPT-5 integration and advanced embedding-based retrieval systems

Design and implement advanced RAG architectures including hybrid search, query routing, and contextual retrieval

Establish multi-agent systems with Agent-to-Agent (A2A) communication protocols and Model Context Protocol (MCP)

Architect Small Language Model (SLM) integration for specialized tasks and cost optimization

Design enterprise-scale solutions supporting millions of documents with sub-second query response times

Establish AI governance, model safety protocols, and regulatory compliance frameworks

Lead architectural reviews for distributed AI systems, microservices, and cloud-native deployments

Hands-on development using Claude Code for rapid prototyping and agentic workflows

Drive architectural reviews for LlamaParse/Azure Document Intelligence integration

Design fault-tolerant, high-availability AI systems with automatic failover and load balancing

Establish comprehensive monitoring, observability, and performance optimization strategies

Mentor technical teams and establish AI engineering best practices using modern toolchains

Oversee model performance evaluation using LangGraph evals and DeepEval frameworks

Mandatory Skills

AI/ML

Azure OpenAI GPT

Agent-to-Agent (A2A) communication

AI & System Architecture

- Define and own the architecture for AI-driven solutions that combine GenAI components (LLMs, RAG, Agentic AI) and ML models (predictive, classification, optimization, forecasting, etc).
- Design scalable, production-ready AI systems that can be embedded within enterprise applications and workflows.
- Establish best practices for MLOps, AIOps, ensuring automation, monitoring, governance, and lifecycle management of solutions.

Generative AI & Machine Learning Integration

- Architect RAG pipelines, multi-agent systems, and hybrid AI architectures that blend LLM reasoning with ML-driven insights.
- Integrate AI solutions with data lakes, knowledge graphs, APIs, and enterprise data platforms.
- Partner with data scientists and ML engineers to optimize ML pipelines, from feature engineering to deployment.

Enterprise System Integration

- Ensure smooth integration of AI systems with CRM, cloud platforms, and existing IT infrastructure.
- Develop secure, API-driven architectures that allow AI solutions to interact seamlessly with business processes.
- Drive architectural decisions for low-latency inference, real-time analytics, and scalable system performance.

Leadership & Collaboration

- Act as a strategic advisor on AI architecture, guiding stakeholders and leadership on opportunities and risks.

- Collaborate with product, engineering, and data teams to align technical solutions with business objectives.
- Mentor engineering teams in AI/ML solution design, best practices, and architectural patterns.

Qualifications & Expertise

- Bachelor's or Master's degree in Computer Science, AI/ML, Data Science, or a related field.
 - 2-3 years of experience in AI engineering, with a focus on applied Generative AI specifically 2+ years of professional experience with NLP, 1+years of professional experience with Large Language Models (LLM)/GenAI technology like OpenAI API, ChatGPT, GPT-4, Bard, Synthesia, Langchain, HuggingFace Transformers, PyTorch and similar, 1+ years of experience with prompt engineering, 1+ years of experience with vector database
 - 2-3 years of experience integrating ML models with user-facing applications.
 - 5-7 years of software architecture, system design, and engineering experience
 - Proven track record in architecting and delivering enterprise-grade AI/ML systems.
 - Strong expertise in LLMs, RAG architectures, multi-agent AI systems, and ML models.
 - Deep understanding of enterprise system integration and API-driven architectures.
 - Proficiency with cloud platforms (AWS, Azure, or GCP), MLOps frameworks, and AI orchestration tools.
 - Hands-on experience with vector databases (Pinecone, Weaviate, FAISS) and retrieval systems.
 - Strong knowledge of Python, PyTorch, TensorFlow, LangChain, or similar AI frameworks.
 - Familiarity with distributed systems, security, and compliance for AI deployments.
 - Experience with speech-to-text / text-to-speech systems for voice interactions.
 - Familiarity with enterprise data integration, security, and compliance considerations.
 - Familiarity with reinforcement learning, reasoning agents, and orchestration frameworks.
 - Experience integrating AI into mission-critical enterprise workflows.
 - Knowledge of AI governance, responsible AI practices, and regulatory compliance.
-
- Partner with a cross-functional team of engineers, research scientists, technical program managers, and product managers to deliver AI-powered products that change how our associates work and how our customers interact with Capital One.
 - Design, develop, test, deploy, and support AI software components including foundation model training, large language model inference, similarity search, guardrails, model evaluation, experimentation, governance, and observability, etc.
 - Leverage a broad stack of Open Source and SaaS AI technologies such as AWS Ultracusters, Huggingface, VectorDBs, Nemo Guardrails, PyTorch, and more.
 - Invent and introduce state-of-the-art LLM optimization techniques to improve the performance — scalability, cost, latency, throughput — of large scale production AI systems.
 - Contribute to the technical vision and the long term roadmap of foundational AI systems at Capital One.

- You love to build systems, take pride in the quality of your work, and also share our passion to do the right thing. You want to work on problems that will help change banking for good.
- Passion for staying abreast of the latest research, and an ability to intuitively understand scientific publications and judiciously apply novel techniques in production.
- You adapt quickly and thrive on bringing clarity to big, undefined problems. You love asking questions and digging deep to uncover the root of problems and can articulate your findings concisely with clarity. You have the courage to share new ideas even when they are unproven.
- You are deeply Technical. You possess a strong foundation in engineering and mathematics, and your expertise in hardware, software, and AI enable you to see and exploit optimization opportunities that others miss.
- You are a resilient trail blazer who can forge new paths to achieve business goals when the route is unknown.

- Architect the Future Define and drive the AI and platform architecture vision, establishing best practices for building secure, scalable, and high-performance systems on Azure, AWS and GCP.
- Lead AI Enablement Design and implement AI/ML capabilities across the platform, including data ingestion, model lifecycle management, and real-time monitoring to deliver actionable insights.
- Evaluate Emerging Technologies Continuously assess new AI/ML tools and trends to ensure the platform remains innovative and aligned with strategic business goals.
- Collaborate Across Functions Act as a key liaison between engineering, product, and business stakeholders, translating business needs into actionable architectural designs and ensuring cohesive execution.
- Establish Technical Standards Define and enforce architectural standards, toolsets, and integration frameworks.
- Optimize Performance and Cost Direct initiatives to monitor and enhance platform performance and cost-efficiency using advanced platform tools, ensuring high availability and scalability.
- Mentor and Inspire Guide and support engineering teams, fostering a culture of innovation, collaboration, and continuous learning within the architecture function.
- Implementing guardrails, optimizing for costs, and evaluating model performances.

- Bachelor's degree in Computer Science, Linguistics, Philosophy, Cognitive Science, or comparable experience in industry or a related field
- 2+ years working with LLMs (e.g., GPT, Claude, Mistral, etc.) in production settings
- Experience with prompt engineering, RAG, or agentic AI architectures
- Strong understanding of conversational UX and prompt evaluation metrics
- Experience with agentic frameworks in practice (langchain, n8n, pydantic, etc)

- Ability to self direct and build full stack prototypes to show thought process and work direction
 - Architect and Design: Lead the design and development of enterprise-grade AI applications and platforms, with a focus on scaling AI solutions for production. This includes defining the technical architecture, selecting appropriate technologies, and ensuring solutions are robust, scalable, and secure.
 - LLM and AI Integration: Integrate and fine-tune Large Language Models (LLMs) and other AI/ML models into enterprise applications. Develop and implement strategies for model deployment, inference, and monitoring, with an emphasis on production-level performance and reliability.
 - Enterprise Architecture: Collaborate with enterprise architects to ensure AI solutions align with the broader company's technical strategy, governance, and standards.
 - Cloud and GenAI Native Development: Design and deploy applications using Cloud Native principles on a hyperscaler platform (AWS, Azure, GCP). Leverage a wide range of hyperscaler tools and services, including containers (Docker, Kubernetes), serverless functions, and managed databases. Should have experience in leveraging various GenAI tools to accelerate software development life cycle.
 - Security & Governance: Ensure the security of all AI/ML systems by addressing potential vulnerabilities such as data privacy concerns, model poisoning, and adversarial attacks.
 - Design Patterns: Apply and enforce Application Design Patterns and Agentic Design Patterns to build resilient and maintainable software systems.
-
- Bachelor's degree in Computer Science or a related technical field.
 - 8+ years' experience as a Software or Solution Architect, with a strong focus on application development and scaling solutions for production environments.
 - Hands-on experience in Java, Python, or .NET.
 - Deep understanding of AI/ML concepts, including experience with LLMs and their application in enterprise settings.
 - Experience implementing multiple AI solutions in a professional, real-world environment.
 - Strong understanding of security implications related to AI/ML systems (e.g., data privacy, model poisoning, adversarial attacks).
 - Familiarity with various hyperscaler tools and services.
 - Hyperscaler Architect certification is required (e.g., AWS Certified Solutions Architect, Azure Solutions Architect Expert, or GCP Professional Cloud Architect).
 - Ability to travel up to 50% based on the work you do and the clients and industries/sectors you serve.
 - Limited immigration sponsorship may be available.
-
- Research, design, development, and deployment of advanced AI agents and agentic systems.

- Architect and implement complex multi-agent systems, including planning, decision-making, and execution capabilities.
- Develop and integrate large language models (LLMs) and other state-of-the-art AI techniques to enhance agent autonomy and intelligence.
- Build robust, scalable, and reliable infrastructure to support the deployment and operation of AI agents at scale.
- Collaborate with product managers, UX designers, and other engineers to define requirements and deliver impactful solutions.
- Diagnose and troubleshoot issues in complex distributed environments and optimize system performance.
- Contribute to the team's technical growth and knowledge sharing.
- Stay up-to-date with the latest advancements in AI research and agentic AI and apply them to our products.
- Leverage enterprise data, market data, and user interactions to build intelligent and personalized agent experiences.
- Contribute to the development of Copilot GenAI Workflows for Users, enabling chat-like command execution.

- Knowledge and passion in machine learning algorithms, Gen AI, LLMs, and natural language processing (NLP).
- Understanding of agent-based modeling, reinforcement learning, and autonomous systems.
- Experience with large language models (LLMs) and their applications in Agentic AI.
- Proficiency in programming languages such as Python, and experience with machine learning frameworks like TensorFlow or PyTorch.
- Experience with cloud platforms (AWS) and containerization technologies (Docker, Kubernetes).
- Understanding of distributed system design patterns and microservices architecture.
- Demonstrable Coding and Algorithms skills.
- Excellent problem-solving and data analysis skills.
- Strong communication and collaboration skills.
- Master's or Ph.D. in Computer Science, Artificial Intelligence, or a related field, or equivalent years of experience.
- Min 2-4+ years of relevant machine learning work experience.

- Architect multi-agent AI solutions to ensure scalability, performance, and seamless integration with RLDatix products.
- Evaluate and select frameworks, tools, and technologies to optimize cost, performance, and maintainability.
- Design, build, and deploy multi-agent AI systems with capabilities in planning, reasoning, memory, and action execution.
- Develop a multi-year AI/ML capability roadmap in collaboration with product, engineering, and data teams.

- Research and innovate by staying ahead of emerging AI technologies and driving practical adoption into production.
- Mentor engineers and foster a culture of technical excellence and innovation.
- Implement evaluation, monitoring, and retraining processes to maintain system accuracy, reliability, and performance.
- Ensure strong governance, security, and ethical AI practices throughout the agent lifecycle.

What Kind of Things We're Most Interested in You Having

- Hands-on expertise in designing and developing multi-agent AI systems.
 - Proven experience creating AI system architectures and roadmaps in production environments.
 - Experience deploying scalable AI/ML applications on cloud platforms (AWS, Azure, or GCP) with MLOps pipelines.
 - Strong programming skills in Python with familiarity across AI/ML libraries and frameworks.
 - In-depth knowledge of Large Language Models (LLMs), orchestration engines, vector databases, and RAG approaches.
 - Strong grasp of software architecture principles, design patterns, and system integration best practices.
 - Background in integrating AI into SaaS platforms (preferred).
 - Deep understanding of data engineering, governance, and privacy regulations in healthcare or other regulated industries.
 - Experience designing and building RESTful APIs and microservices for AI/ML integration.
 - Excellent communication skills with the ability to bridge technical and business perspectives.
 - A collaborative, problem-solving mindset with strategic thinking to design complex systems.
-
- Research, design, development, and deployment of advanced AI agents and agentic systems.
 - Architect and implement complex multi-agent systems, including planning, decision-making, and execution capabilities.
 - Develop and integrate large language models (LLMs) and other state-of-the-art AI techniques to enhance agent autonomy and intelligence.
 - Build robust, scalable, and reliable infrastructure to support the deployment and operation of AI agents at scale.
 - Collaborate with product managers, UX designers, and other engineers to define requirements and deliver impactful solutions.
 - Diagnose and troubleshoot issues in complex distributed environments and optimize system performance.
 - Contribute to the team's technical growth and knowledge sharing.
 - Stay up-to-date with the latest advancements in AI research and agentic AI and apply them to our products.

- Leverage enterprise data, market data, and user interactions to build intelligent and personalized agent experiences.
- Contribute to the development of Copilot GenAI Workflows for Users, enabling chat-like command execution.

Qualifications

- Knowledge and passion in machine learning algorithms, Gen AI, LLMs, and natural language processing (NLP).
 - Understanding of agent-based modeling, reinforcement learning, and autonomous systems.
 - Experience with large language models (LLMs) and their applications in Agentic AI.
 - Proficiency in programming languages such as Python, and experience with machine learning frameworks like TensorFlow or PyTorch.
 - Experience with cloud platforms (AWS) and containerization technologies (Docker, Kubernetes).
 - Understanding of distributed system design patterns and microservices architecture.
 - Demonstrable Coding and Algorithms skills.
 - Excellent problem-solving and data analysis skills.
 - Strong communication and collaboration skills.
 - Master's or Ph.D. in Computer Science, Artificial Intelligence, or a related field, or equivalent years of experience.
 - Min 1-3+ years of relevant machine learning work experience.
-
- Design and develop Generative AI solutions using cloud-based managed AI services (e.g. Amazon Bedrock, Amazon SageMaker, Google Cloud Vertex AI).
 - Design and develop AI Agent workflows using cloud and open source agentic frameworks.
 - Containerize AI applications and deploy them using cloud orchestration services.
 - Collaborate with data architects/engineers to build end-to-end AI pipelines.
 - Implement MLOps practices to automate the development, deployment, and monitoring of AI applications and models.
 - Implement and manage robust monitoring systems for AI solutions in production environments, ensuring continuous performance tracking, anomaly detection, and model drift analysis; collaborate with cross-functional teams to deploy model updates, maintain version control, and optimize model efficiency over time.
 - Use Infrastructure as Code (IaC) to manage and version cloud resources for AI projects.
 - Ensure clear and accessible knowledge transfer to internal teams and create knowledge-sharing resources to ensure smooth transitions during model handoffs and system updates.
 - Contribute to the development of best practices and standards for AI engineering within the organization.

Qualifications

Minimum Requirements

- Bachelor's degree and 6 years' experience in Architect/Engineer roles OR a Master's degree and 4 years' experience Architect/Engineer roles OR a Ph.D. and 1 years' experience in Architect/Engineer roles OR 10 years' experience in Architect/Engineer roles.
- Experience using Python programming language to create and integrate GenAI solutions related to business workloads
- 2 years of experience of managing AI services within one cloud platform (AWS preferred, Google Cloud, Azure)
- Experience with container services and orchestration (AWS EKS, ECS, GKE, etc.)
- Experience in common machine learning, deep learning, and LLM frameworks, such as TensorFlow, PyTorch, scikit-learn, Hugging Face Transformers, LangChain, LangGraph.
- 2 years experience with using Terraform for managing Infrastructure as Code (IaC) as well as CI/CD deployments.

AI/ML Strategy & Solution Architecture:

- Research, design, and prototype state-of-the-art AI/ML and Generative AI models (e.g., LLMs, Transformers, Diffusion Models) to address complex business challenges.
- Architect end-to-end ML pipelines, encompassing data ingestion, preprocessing, model training, evaluation, deployment, and monitoring (MLOps).
- Design cloud-native, scalable, and cost-optimized AI solutions for deployment on AWS (SageMaker, Bedrock) and GCP (Vertex AI, Gemini).
- Make strategic build-vs-buy decisions and select the right frameworks (e.g., TensorFlow, PyTorch, Hugging Face, LangChain).
- Use experience with Multi Model RAG based solution development & Agentic AI based solution approaches.

Hands-on Development & Leadership:

- Lead by example: write production-quality code, build and tune models, and troubleshoot complex issues within the data and ML stack.
- Own the entire ML lifecycle, ensuring best practices in reproducibility, versioning (e.g., MLflow, DVC), and model governance.
- Collaborate closely with the Cloud/DevOps Architect to integrate AI workloads seamlessly into CI/CD pipelines and cloud infrastructure.

Mentorship & Team Development:

- Serve as the primary technical mentor for interns and junior data scientists, providing guidance on projects, code reviews, and research methodologies.
- Foster a culture of continuous learning by conducting workshops on advanced AI topics, ethical AI, and new technologies.
- Translate complex AI concepts into actionable tasks for the team.

Pre-Sales & Market Evangelism:

- Act as the key AI expert during pre-sales, demonstrating the technical superiority and value of our AI solutions to potential customers and partners.

- Architect and deliver compelling proof-of-concepts (POCs) and demos that showcase practical applications.
- Create high-quality technical content (whitepapers, architecture diagrams, blog posts) to articulate the innovation behind our marketplace offerings.
- Present solutions confidently to both technical and non-technical audiences, supporting the sales cycle.

Required Experience and Education:

- 5-10 years of hands-on experience in data science, machine learning, and AI, with at least 3 years in an architect or tech lead role.
 - Bachelor's Degree preferred.
 - A proven track record of building, training, tuning, and deploying machine learning models into production environments.
 - Deep, practical experience with Generative AI, including working with Large Language Models (LLMs), prompt engineering, RAG architectures, and fine-tuning.
 - Strong proficiency in Python and its core ML ecosystem (e.g., Pandas, NumPy, Scikit-learn, TensorFlow/PyTorch, matplotlib). Scala, GO experience is an added advantage.
 - Experience in building AI solutions which follow, "secure by design" principles, including data encryption, privacy-preserving ML techniques, vulnerability scanning, and compliance readiness (SOC2, FedRAMP).
 - Extensive experience with cloud AI platforms: AWS SageMaker and/or GCP Vertex AI.
 - Solid understanding of MLOps principles and tools (e.g., MLflow, Kubeflow).
 - Experience with data engineering and processing large datasets (SQL, Spark, Leetcode, Strata Scratch).
 - A proven ability to mentor and develop junior talent.
-
- Serve as an expert in implementing and deploying AI applications based on large language models (LLMs), internal and external agentic frameworks, and custom models.
 - Work with hardware architects to identify how to best design, customize, and deploy AI-based solutions to their specific problem domains.
 - Collaborate with infrastructure engineers to improve existing automated workflows by incorporating LLMs and establishing best practices for future solutions.
 - Develop and optimize retrieval and generation algorithms for enterprise data (text, code, and images) to build advanced AI applications.
 - Interact with internal research groups on how to solve complex chip design problems in new ways by leveraging machine learning (ML) and deep learning (DL).
 - Research emerging AI technologies and engineering best practices to continuously evolve our development ecosystem and maintain a competitive edge.

What We Need To See

- MSc or PhD in Data Science, Computer Science/Engineering, Electrical Engineering, or equivalent experience.
- 5+ years of industry or research experience.

- Deep practical knowledge of LLMs, DL/ML, and Agent development.
 - Well versed in agentic literature and eager to continue learning.
 - Strong background in implementing AI solutions to solve real-world engineering problems.
 - Experience with training/fine-tuning custom models, building multi-agent systems, retrieval augmented generation (RAG) pipelines, and vector databases.
 - Strong analytical, communication, and interpersonal skills.
- Design comprehensive AI solution and technology architecture, integrating latest AI technology developments into world-class solutions.
 - Lead high-level architectural discussions with clients, providing expert guidance on best practices for AI implementations across AI PC, Edge, Data Centre and Public Cloud environments.
 - Ensure solutions align with modern best practices across the full spectrum of platforms and environments. Deep understanding across GPU/NPU, Cognitive Infrastructure, Application and Copilot/agent domains.
 - Contribute to HCL's thought leadership in the AI & Cloud domains with a deep understanding of open[1]source technologies (e.g., Kubernetes, OPEA) and partner technologies.
 - Collaborate on joint technical projects with global partners, including Google, Microsoft, AWS, NVIDIA, IBM, Red Hat, Intel, and Dell

Mandatory Skills & Experience

- 10+ years architecture design
- 7+ years in software engineering Technologies:
- Professional-level expertise in Public Cloud environments (AWS, Azure, Google Cloud).
- Demonstrable coding proficiency with Python, Java or Go languages.
- AI Expertise:
- Advanced machine learning algorithms, GenAI models (e.g., GPT, BERT, DALL-E, GEMINI), NLP techniques

Desired Skills & Experience

- Knowledge of GenAI operations (LLMOps), experience Governing AI models in production environments.
- Proficiency in data engineering for AI, including data preprocessing, feature engineering, and pipeline creation.
- Expertise in AI model fine-tuning and evaluation, with a focus on improving performance for specialized tasks.
- Copilot design, engineering and extensions.
- Knowledgeable about Responsible AI, including governance and ethics. Bias mitigation, with experience in implementing strategies to ensure fair and unbiased AI solutions.
- Deep Learning Frameworks (TensorFlow, PyTorch)

- Innovation and Emerging Technology Trends
- Strategic AI Vision and Road mapping
- Enthusiastic about working in a fast-paced environment using the latest technologies, and passionate about HCL's dynamic and high-energy Lab culture.
 - Technical Vision
 - Shape the long-term technical strategy for Generative AI at DataRobot, ensuring our systems are scalable, maintainable, and aligned with business goals.
 - Lead architectural decisions for GenAI tooling (e.g. Agents and Agentic Workflows, Prompt Management, Frameworks & Libraries, LLM Onboarding, Tools, Multi-Agent Evaluations, Multimodality, etc.) and GenAI systems (e.g. Inference optimization, Distributed Training, Finetuning, Distillation, LLM Moderation and Guardrails, etc.).
 - Anticipate technical risks and propose mitigation strategies before they become roadblocks.
 - Execution & Impact
 - Hands-on development: Build, ship and operate critical AI infrastructure, balancing rapid iteration with long-term technical health.
 - Solve cross-pillar challenges, such as optimizing LLM latency/cost, improving Agentic workflows, RAG workflows, or ensuring reproducibility in generative outputs.
 - Drive operational excellence by improving observability, reliability, and performance of AI systems in production.
 - Leadership & Mentorship
 - Act as the technical advisor to engineering leadership (VPs, Directors) and product teams, influencing roadmap priorities.
 - Mentor Staff/Senior Engineers, elevating the team's technical bar through design reviews, best practices, and knowledge sharing.
 - Foster a culture of innovation and rigor, ensuring AI solutions are both cutting-edge and production-ready.
 - Cross-Functional Collaboration
 - Partner with Research, Product, and GTM teams to align technical efforts with customer needs and market opportunities.
 - Communicate complex technical concepts to executives and non-technical stakeholders, enabling data-driven decisions.

Knowledge, Skills & Abilities:

- A technical leader with 8+ years of software engineering experience, including 3+ years in AI/ML systems (Generative AI preferred).
- Deep expertise in:
 - Scalable systems: Distributed computing, containerization (Docker/K8s), and multi-cloud deployments (AWS/GCP/Azure/Openshift).
 - Software craftsmanship: Python, API design, and infrastructure-as-code (Terraform, Pulumi).

- Proven track record of delivering multi-year, high-impact projects with cross-functional dependencies.
 - A systems thinker who balances innovation with operational pragmatism (e.g., cost-to-serve, technical debt).
 - Exceptional communication and influence skills—you can rally teams behind a technical vision.
-
- Deep expertise in Generative AI stack: LLMs (open/closed-source), fine-tuning, RAG, orchestration, and evaluation metrics.
 - Experience with MLOps tools and AI-specific infra (e.g., vector DBs, GPU optimization).
 - Contributions to open-source AI projects or published research in ML/GenAI.
 - Background in enterprise SaaS or B2B AI products
-
- Architect and implement scalable agent orchestration patterns (graph-based executors, state management, multi-agent coordination) for production workloads.
 - Own critical integrations: model adapters, LLM gateway hooks, vector DBs, tools & external APIs, and the platform's LLMops flows.
 - Build and improve tracing, benchmarking and observability for LLMs and agents — token/cost accounting, latency p95, throughput, and correctness checks.
 - Drive design for safety/guardrails: moderation hooks, human-in-the-loop checkpoints, replayable audit trails and policy enforcement.
 - Mentor junior engineers, run design reviews, and improve engineering practices (testing, CI/CD, chaos testing for agents).
 - Work directly with strategic customers to prototype complex agentic solutions and translate them into product features.

Must-have

- 5–8 years of software engineering with substantial experience building distributed systems, infra, or ML platforms.
- Deep practical experience integrating and deploying LLMs in production (RAG, retrieval, embeddings pipelines).
- Hands-on experience with agent orchestration frameworks (LangGraph / LangChain or custom agent runtimes) and stateful workflow design.
- Strong systems knowledge: Kubernetes, container orchestration, service meshes, and performance tuning.
- Proven track record building observability, cost controls, and policy enforcement for production services.

Preferred / differentiators

- Experience building or contributing to open-source LLM orchestration tools (LangGraph, LangChain, or similar).
- Familiarity with enterprise constraints: on-prem/cloud hybrid deployments, data residency, compliance requirements.
- Background in security, privacy, or model governance for LLMs.
- Demonstrated leadership in cross-functional projects and direct customer engagement.

Generative AI Strategy & Architecture

- Define and evolve the enterprise-wide architectural vision and strategy for Generative AI and agentic systems, aligning with overall business goals and technology roadmaps.
- Lead the design of modular, reusable, and scalable architectural patterns for GenAI and agentic applications across various domains.
- Design and implement robust, secure solution patterns that can be operationalized across enterprise environments.

Foundation Models Expertise

- Deep expertise in various foundation models (e.g., LLMs, vision models, multimodal models) including their architectures, strengths, limitations, and fine-tuning techniques.
- Evaluate, select, and integrate appropriate foundation models for specific use cases, considering factors like performance, cost, and ethical implications.
- Develop strategies for model pre-training, fine-tuning, and continuous improvement.

Agentic AI Systems

- Architect and implement intelligent agentic AI systems using frameworks like LangChain, LangGraph, CrewAI, AutoGen, or similar.
- Design complex agentic workflows, including planning, reasoning, tool integration, memory management, and human-in-the-loop interactions.
- Lead the integration of agentic AI solutions with existing enterprise systems and define integration standards (e.g., RESTful APIs, microservices).

Prompt Engineering & Orchestration

- Apply advanced prompt engineering techniques to optimize AI model performance and steer outputs towards desired outcomes, minimizing bias and hallucinations.
- Oversee orchestration of AI components and services, including LLM APIs, vector databases, and external tools.
- Develop and implement Retrieval Augmented Generation (RAG) based solutions for enhanced contextual understanding and factual accuracy.

Scalable Infrastructure & MLOps

- Design and build cloud-native, containerized infrastructure (e.g., Kubernetes, ECS, EKS on AWS, Azure, GCP) for deploying GenAI models and agentic systems at scale.
- Implement robust MLOps pipelines for the continuous integration, delivery, deployment, monitoring, and management of AI models in production.
- Ensure AI solutions comply with regulatory requirements, data privacy, and ethical AI standards.

Innovation & Thought Leadership

- Stay abreast of the latest advancements in AI, machine learning, deep learning, and agentic systems, and apply this knowledge to drive innovation.
 - Rapidly develop Proofs of Concept (PoCs) and iterate solutions using an agile, experimental approach without compromising architectural integrity or long-term scalability.
 - Serve as a technical expert and mentor to junior AI engineers and architects, fostering a culture of continuous learning and improvement.
 - Contribute to internal workshops, knowledge sharing, and external forums as a thought leader.
-
- Collaborate effectively with cross-functional teams including product managers, data scientists, software engineers, security, and business stakeholders.
 - Articulate complex technical concepts to diverse audiences, both technical and non-technical.
 - Maintain detailed architectural documentation and operational playbooks.
-
- Bachelor's or Master's degree in Computer Science, Artificial Intelligence, Machine Learning, or a related quantitative field.
 - 10+ years of progressive experience in software architecture, with at least 2-3 years specifically focused on Generative AI and Machine Learning.
 - Demonstrated expertise in designing and implementing AI architectures with a strong focus on Generative AI and Agentic AI technologies.
 - Profound understanding of various foundation models (LLMs, vision models, multimodal models) and their underlying architectures (e.g., Transformers).
 - Hands-on experience with agentic AI frameworks such as LangChain, LangGraph, CrewAI, AutoGen, or similar.
 - Strong programming skills in Python is essential, with experience in relevant AI/ML frameworks (e.g., TensorFlow, PyTorch).
 - Experience with cloud platforms (AWS, Azure, GCP) and their AI/ML services (e.g., Amazon SageMaker, Azure Machine Learning, Google Cloud AI Platform, Vertex AI).
 - Familiarity with MLOps principles and tools for deploying, monitoring, and managing AI models and agentic systems in production (e.g., MLflow, Kubeflow).

- Experience with vector databases (e.g., Pinecone, Weaviate, FAISS, Azure AI Search) and techniques for processing and ingesting unstructured data.
- Excellent communication, interpersonal, and leadership skills, with the ability to influence and drive technical decisions.
- Strong analytical and problem-solving abilities.
- Willingness to work in the New York office 3 days per week.
- Hands-on builder who ships—you're comfortable owning problems end-to-end in a fast, early-stage environment.
- Deep experience with LLMs, agent architectures, and turning research into reliable products.
- Strong product sense: you translate messy legal/medical workflows into usable, trustworthy automation.
- High ownership and bias to action; you thrive with autonomy and ambiguity.
- Customer-obsessed, rapidly learns new domains and iterates based on impact.

What You'll Do

- Design and ship autonomous agentic systems that plan, reason, and execute multi-step workflows.
- Build at scale: ingest and structure millions of unstructured documents (medical records, insurance files, legal PDFs).
- Fine-tune and deploy language models for extraction, entity linking, retrieval, and predictive analytics.
- Own the stack—from modeling and tool orchestration to production deployment, monitoring, and iteration.
- Define evaluation/measurement frameworks; improve quality with experiments and data analysis.
- Partner with legal/clinical experts and engineering to deliver measurable business outcomes.

Qualifications

- Proven track record in building and deploying production ML/LLM systems (agent frameworks, tool use, orchestration).
- Strong with Python and modern ML tooling (vector stores, embeddings/RAG, evaluation, tracing); solid software engineering fundamentals.
- Experience processing unstructured documents at scale (pipelines, OCR, normalization, metadata).
- Cloud/MLOps experience (AWS/GCP/Azure, containers, CI/CD, observability).
- Excellent judgment, communication, and bias-for-evidence; write clear, discoverable notes and docs.
- Nice to haves: legal/healthcare data exposure; safety/ethics/governance for AI; startup experience.

- Design and implement RAG-based solutions to enhance LLM capabilities with external knowledge sources
- Develop and optimize LLM fine-tuning strategies for specific use cases and domain adaptation
- Create robust evaluation frameworks for measuring and improving model performance
- Build and maintain agentic workflows for autonomous AI systems
- Collaborate with cross-functional teams to identify opportunities and implement AI solutions

Minimum Qualifications

- Bachelor's degree in Computer Engineering, Computer Science, Electrical Engineering, or related field and 4+ years of Software Engineering, Electrical Engineering, or related work experience.

OR

Master's degree in Computer Engineering, Computer Science, Electrical Engineering, or related field and 3+ years of Software Engineering, Electrical Engineering, or related work experience.

OR

PhD in Computer Engineering, Computer Science, Electrical Engineering, or related field and 2+ years of Software Engineering, Electrical Engineering, or related work experience.

Required Qualifications

- Bachelor's or Master's degree in Computer Science, or related technical field
- 3+ years of experience in Machine Learning/AI engineering
- Strong programming skills in Python and experience with ML frameworks (PyTorch, TensorFlow)
- Practical experience with LLM deployments and fine-tuning
- Experience with vector databases and embedding models
- Familiarity with modern AI/ML infrastructure and cloud platforms (AWS, GCP, Azure)
- Strong understanding of RAG architectures and implementation

Preferred Qualifications

- Experience with popular LLM frameworks (Langchain, LlamaIndex, Transformers)
- Knowledge of prompt engineering and chain-of-thought techniques

- Experience with containerization and microservices architecture
 - Background in Reinforcement Learning
 - Contributions to open-source AI projects
 - Experience with ML ops and model deployment pipelines
-
- Own the end-to-end architecture for Zendesk's GenAI platform, ensuring alignment with business goals and technical best practices.
 - Set technical direction for core systems including LLM Proxy, agent orchestration layers, evaluation and benchmarking pipelines, and model observability tooling
 - Design and implement model routing, fallback strategies, and A/B testing infrastructure for LLMs from multiple vendors
 - Establish engineering standards for safety, latency, cost attribution, and reliability across all GenAI services.
 - Coach Staff and Senior Engineers, providing deep technical guidance and fostering a culture of technical excellence.

What You Bring To The Role

- 10+ years in ML/AI engineering, including at least 3 years in staff/principal-level platform or infrastructure leadership roles..
- Expertise in LLM systems, multi-model orchestration, and GenAI infrastructure patterns
- Proven ability to design complex, distributed systems with high reliability, security, and scalability requirements
- Strong experience with AWS, GCP, or Azure; Kubernetes; Docker; and distributed event-driven architectures
- Fluency in Python and at least one other server-side language (Java, Scala, Golang, or Ruby).
- Track record of driving alignment across multiple engineering teams and influencing product direction.

Preferred Qualifications

- Background in agentic architectures and complex workflow orchestration for AI agents.
- Experience and contributions to enterprise-scale ML platforms.

What Our Tech Stack Looks Like

- Our code is written in Python.
- Our servers live in AWS.
- LLM Vendors: OpenAI, Anthropic, Google, Llama
- Infra: Kubernetes, Docker, Kafka, AWS

onfido is building an AI-native operating system for brands that sell into brick-and-mortar retail. We're building intuitive and powerful software paired with AI agents that help brands automate the

1000s of hours of tedious manual work that goes into scaling a brand in retail. Brands like Tropicana and Olipop use our first products in accounting and planning to account for and plan for \$Bs of retail revenue.

We have best-in-class growth (5x YoY), efficiency, and we've recently raised a \$15M Series A to build on the existing momentum.

You will be our Founding Staff ML / AI Engineer, setting the foundation for evolving our AI / ML capabilities. We have a unique mix of work across LLMs and advanced machine learning.

- Developing agentic systems to retrieve information from legacy systems
 - Leverage NLP and LLMs to understand complex documents
 - Predict brand sales using advanced machine learning techniques
 - Detect anomalies in sales data to help brands sell more
 - Recommend the optimal promotion for a brand at a retailer
-
- Own the full AI / ML lifecycle—from research & prototyping to production deployment & monitoring
 - Collaborate directly with our CTO and CEO to define our technical and product strategy.
-
- Minimum of 5 years of Applied AI / ML experience, ideally in a startup environment
 - Strong product mindset—the ability to bridge AI capabilities with real-world usage
 - Experience fine-tuning LLMs (Llama, GPT) and designing AI prompt workflows.
 - Strong background in ML systems design and engineering best practices
-
- Architect the Future Define and drive the AI and platform architecture vision, establishing best practices for building secure, scalable, and high-performance systems on Azure.
-
- Lead AI Enablement Design and implement AI/ML capabilities across the platform, including data ingestion, model lifecycle management, and real-time monitoring to deliver actionable insights.
-
- Evaluate Emerging Technologies Continuously assess new AI/ML tools and trends to ensure the platform remains innovative and aligned with strategic business goals.
-
- Collaborate Across Functions Act as a key liaison between engineering, product, and business stakeholders, translating business needs into actionable architectural designs and ensuring cohesive execution.
-
- Establish Technical Standards Define and enforce architectural standards, toolsets, and integration frameworks which will deliver on Intapp's promise. Lead governance boards to

review and approve major technical decisions.

- Optimize Performance and Cost Direct initiatives to monitor and enhance platform performance and cost-efficiency using advanced Azure tools, ensuring high availability and scalability.
- Mentor and Inspire Guide and support engineering teams, fostering a culture of innovation, collaboration, and continuous learning within the architecture function.
- Experience 10+ years in enterprise platform architecture, with expertise in public cloud deployments like Azure, AWS, GCP.
- Leadership & Communication Proven ability to lead cross-functional teams, influence strategic direction, and communicate complex technical concepts to diverse audiences.
- AI/ML Expertise Strong background in AI/ML frameworks, modern tool chains, ability to advocate for generative and discriminative models to deliver AI capabilities at Enterprise Scale.
- Software Engineering Skills Proficiency in modern programming languages, SQL and NoSQL databases, and event-driven systems like Kafka, with a focus on building scalable, maintainable solutions.
- Security & Scalability Focus Demonstrated success in designing secure, scalable platforms with a commitment to operational excellence and resource optimization.
- Hands-on experience with MLOps, model management, and AI observability tools
- Familiarity with containerization and orchestration (e.g., Docker, Kubernetes)
- Knowledge of Responsible AI practices and data governance frameworks
- Passion for staying ahead of industry trends in AI and platform architecture
- Advanced degree in Computer Science, Engineering, or a related field (or equivalent experience)

- Drive innovation by experimenting with new methodologies and harnessing the power of LLMs to develop groundbreaking features for our AI assistants
- Lead ideas from concept to reality—crafting prototypes, engineering robust code, and ensuring impactful deployment
- Collaborate closely with visionary engineers, product managers, technical program managers, and other key stakeholders across North America and Europe
- Tackle real business and user challenges from various industries—then architect scalable, generalized solutions that will shape the foundation of our AI systems

We Would Love To Hear From You If

- You have at least 8 years of professional experience as a Software Engineer with a Bachelor's in Computer Science or relevant field, or at least 6 years of professional experience with a Master's degree.
 - You have deep expertise in AI/ML or a closely related discipline, with a proven record of delivering impactful solutions.
 - You bring hands-on experience with a variety of agentic frameworks and techniques such as LangGraph, RAG, DSPy, and BAML in implementing AI-driven systems within diverse, global teams.
-
- Develop and fine-tune Large Language Models (LLMs) and foundational models to meet specific product and business goals
 - Design and implement scalable machine learning pipelines for training, fine-tuning, and deploying LLMs
 - Build and optimize agentic workflows, enabling LLMs to interface with external tools, APIs, and environments
 - Collaborate with cross-functional teams including data scientists, engineers, and product leaders to embed AI capabilities into FOX applications
 - Continuously evaluate and improve model performance, reliability, and efficiency in real-world use cases
 - Stay up to date with the latest research in Generative AI, autonomous agents, and multi-modal learning
 - Conduct rigorous testing, validation, and monitoring of deployed models in production environments
 - Develop best practices, documentation, and internal tooling to support ongoing innovation and scalability
 - Mentor junior engineers and help foster a high-performing, collaborative team culture
 - Participate in system design and code reviews to uphold strong engineering standards
-
- Proven experience developing and deploying machine learning models in production, with a focus on LLMs or generative AI

- Deep understanding of transformer architectures and techniques for fine-tuning and prompt engineering
 - Strong proficiency in Python and ML frameworks like PyTorch, TensorFlow, or JAX
 - Experience building agentic systems, including tool use, function calling, and multi-step reasoning
 - Familiarity with vector databases, embeddings, and retrieval-augmented generation (RAG) techniques
 - Hands-on experience with cloud platforms and ML infrastructure (e.g., AWS, GCP, Azure, or Databricks)
 - Solid software engineering skills, including version control, CI/CD, and system design
 - Ability to communicate complex technical concepts clearly to both technical and non-technical audiences
 - A collaborative mindset and ability to thrive in a fast-paced, cross-functional environment
-
- Experience with autonomous agents, LangChain, OpenAI functions, or similar orchestration frameworks
 - Familiarity with ethical and safety considerations in AI development
 - Prior experience mentoring other engineers or leading projects
 - Experience building multi-modal models (e.g., combining text, vision, or audio)
 - Contributions to open-source ML libraries or published research in relevant areas